

Sample API Callouts — AI Email Reply Generator

Sample API Callouts — AI Email Reply Generator

Company: MDS LLC

Package: AI Email Reply Generator (namespace: aiemailreply)

Date: May 7, 2026

Overview

This document provides sample HTTP requests and responses for all external callouts made by the AI Email Reply Generator. All callouts originate from within the **customer's Salesforce org** and are directed to the **customer's own AI provider account** using the customer's own API key. MDS LLC receives no data.

All callouts are performed by `AIProviderCalloutService.callProvider()` via Salesforce Named Credentials. The Authorization header is injected automatically by the Named Credential's `generateAuthorizationHeader` mechanism — the API key never appears in Apex code or logs.

Callout trigger: User clicks “Generate Reply” in the Lightning component.

Data sent: Assembled prompt containing the email thread subject, body excerpt, selected tone, and optional custom instructions.

Data received: AI-generated reply text and token usage counts.

Callout 1 — OpenAI

Named Credential: AI_Email_Reply_Provider

Default Endpoint: `https://api.openai.com/v1/chat/completions`

HTTP Request

```
POST /v1/chat/completions HTTP/1.1
Host: api.openai.com
Content-Type: application/json
Authorization: Bearer sk-..... (injected by Named Credential)
Content-Length: 612
```

Request Body:

```
{
  "model": "gpt-4o",
  "temperature": 0.7,
  "max_tokens": 500,
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful email assistant."
    },
    {
      "role": "user",
```

```
"content": "Write a professional email reply to the following message:\n\nSubject: Question about order #12345\n\nHi, I wanted to check on the status of my recent order. It has been a week and I haven't received any updates.\n\nTone: Professional\nInstructions: Be concise and reassuring."
    }
  ]
}
```

HTTP Response

HTTP/1.1 200 OK
Content-Type: application/json
x-request-id: req_abc123def456
openai-processing-ms: 842
openai-version: 2020-10-01
x-ratelimit-limit-requests: 10000
x-ratelimit-remaining-requests: 9998
x-ratelimit-reset-requests: 6ms

Response Body:

```
{
  "id": "chatcmpl-abc123def456xyz",
  "object": "chat.completion",
  "created": 1746633600,
  "model": "gpt-4o-2024-08-06",
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "Dear Customer,\n\nThank you for reaching out regarding order #12345. I apologize for the lack of updates and understand your concern.\n\nI have checked your order status and can confirm it is currently being processed and will ship within the next 1-2 business days. You will receive a tracking number via email once it has been dispatched.\n\nPlease don't hesitate to contact us if you have any further questions.\n\nBest regards,\nSupport Team"
      },
      "finish_reason": "stop"
    }
  ],
  "usage": {
    "prompt_tokens": 98,
    "completion_tokens": 97,
    "total_tokens": 195
  }
}
```

Fields extracted by AIProviderResponseParser: - choices[0].message.content → reply text returned to user - usage.total_tokens → stored on AI_Email_Reply_Draft__c.Token_Count__c

Callout 2 — Anthropic

Named Credential: AI_Email_Reply_Anthropic

Default Endpoint: <https://api.anthropic.com/v1/messages>

HTTP Request

```
POST /v1/messages HTTP/1.1
Host: api.anthropic.com
Content-Type: application/json
Authorization: Bearer sk-ant-..... (injected by Named Credential)
anthropic-version: 2023-06-01
Content-Length: 489
```

Request Body:

```
{
  "model": "claude-opus-4-5",
  "max_tokens": 500,
  "temperature": 0.7,
  "messages": [
    {
      "role": "user",
      "content": "Write a professional email reply to the following message:\n\nSubject:
                  Question about order #12345\n\nHi, I wanted to check on the status of my recent
                  order. It has been a week and I haven't received any updates.\n\nTone:
                  Professional\nInstructions: Be concise and reassuring."
    }
  ]
}
```

Note: Anthropic's API does not use a separate system role message field at the top level in this integration. The system instruction is embedded in the prompt.

Note: The anthropic-version: 2023-06-01 header is set explicitly by AIProviderCalloutService for all Anthropic callouts.

HTTP Response

```
HTTP/1.1 200 OK
Content-Type: application/json
request-id: req_01AbCdEfGhIjKl
anthropic-ratelimit-requests-limit: 1000
anthropic-ratelimit-requests-remaining: 999
anthropic-ratelimit-tokens-limit: 80000
anthropic-ratelimit-tokens-remaining: 79805
```

Response Body:

```
{
  "id": "msg_01AbCdEfGhIjKlMnOpQrStUv",
  "type": "message",
  "role": "assistant",
  "model": "claude-opus-4-5-20251101",
  "content": [
    {
      "type": "text",
      "text": "Dear Customer,\n\nThank you for contacting us about order #12345. I apologize
              for the absence of updates and appreciate your patience.\n\nYour order is currently
              in our fulfillment queue and is expected to ship within 1-2 business days. A
              tracking confirmation will be sent to your registered email address upon
              dispatch.\n\nWe value your business and are committed to resolving this promptly.
              Please reach out if you need any further assistance.\n\nKind regards,\nCustomer
              Support"
    }
  ],
  "stop_reason": "end_turn",
  "usage": {
```

```

    "input_tokens": 95,
    "output_tokens": 110
  }
}

```

Fields extracted by AIPProviderResponseParser: - content[0].text → reply text returned to user -
usage.input_tokens + usage.output_tokens → summed and stored as token count

Callout 3 — Google Gemini

Named Credential: AI_Email_Reply_Google

Default Endpoint: https://generativelanguage.googleapis.com/v1beta/models

Full URL (runtime): https://generativelanguage.googleapis.com/v1beta/models/gemini-2.0-flash:generateContent

The model name and :generateContent action are appended to the base URL at runtime by
AIPProviderCalloutService.buildProviderEndpoint() using EncodingUtil.urlEncode(model, 'UTF-8').

HTTP Request

```

POST /v1beta/models/gemini-2.0-flash:generateContent HTTP/1.1
Host: generativelanguage.googleapis.com
Content-Type: application/json
Authorization: Bearer AIza..... (injected by Named Credential)
Content-Length: 598

```

Request Body:

```

{
  "model": "gemini-2.0-flash",
  "temperature": 0.7,
  "generationConfig": {
    "maxOutputTokens": 500,
    "temperature": 0.7
  },
  "contents": [
    {
      "role": "user",
      "parts": [
        {
          "text": "Write a professional email reply to the following message:\n\nSubject: Question about order #12345\n\nHi, I wanted to check on the status of my recent order. It has been a week and I haven't received any updates.\n\nTone: Professional\nInstructions: Be concise and reassuring."
        }
      ]
    }
  ]
}

```

HTTP Response

```

HTTP/1.1 200 OK
Content-Type: application/json; charset=UTF-8

```

Vary: Origin, X-Origin
 x-request-id: xyz789abc123

Response Body:

```
{
  "candidates": [
    {
      "content": {
        "parts": [
          {
            "text": "Dear Customer,\n\nThank you for reaching out regarding order #12345. I sincerely apologize for the lack of communication regarding your order.\n\nI've checked your order and it is currently being prepared for shipment. You can expect it to ship within the next 1-2 business days, at which point you'll receive an email with your tracking information.\n\nWe appreciate your patience and value your business. Please feel free to reach out if you have any other questions.\n\nSincerely,\nCustomer Support Team"
          }
        ],
        "role": "model"
      },
      "finishReason": "STOP",
      "index": 0
    }
  ],
  "usageMetadata": {
    "promptTokenCount": 88,
    "candidatesTokenCount": 107,
    "totalTokenCount": 195
  },
  "modelVersion": "gemini-2.0-flash"
}
```

Fields extracted by AIPProviderResponseParser: - candidates[0].content.parts[0].text → reply text
 returned to user - usageMetadata.totalTokenCount → stored as token count

Callout 4 — xAI (Grok)

Named Credential: AI_Email_Reply_XAI

Default Endpoint: https://api.x.ai/v1/chat/completions

xAI's API is OpenAI-compatible. The request and response format are identical to Callout 1 (OpenAI). AIPProviderResponseParser uses the parseOpenAIResponse path for xai.

HTTP Request

```
POST /v1/chat/completions HTTP/1.1
Host: api.x.ai
Content-Type: application/json
Authorization: Bearer xai-..... (injected by Named Credential)
Content-Length: 608
```

Request Body:

```
{
  "model": "grok-3",
```

```

"temperature": 0.7,
"max_tokens": 500,
"messages": [
  {
    "role": "system",
    "content": "You are a helpful email assistant."
  },
  {
    "role": "user",
    "content": "Write a professional email reply to the following message:\n\nSubject:
  Question about order #12345\n\nHi, I wanted to check on the status of my recent
  order. It has been a week and I haven't received any updates.\n\nTone:
  Professional\nInstructions: Be concise and reassuring."
  }
]
}

```

HTTP Response

HTTP/1.1 200 OK
 Content-Type: application/json
 x-request-id: xai-req-abc123xyz

Response Body:

```

{
  "id": "chatcmpl-xai-abc123xyz",
  "object": "chat.completion",
  "created": 1746633600,
  "model": "grok-3",
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "Dear Customer,\n\nThank you for your patience and for reaching out about
        order #12345. I apologize for not keeping you better informed about your order's
        progress.\n\nI'm pleased to let you know that your order is currently being prepared
        for shipment and should be dispatched within 1-2 business days. Once shipped, you
        will receive an email confirmation with tracking details.\n\nThank you for your
        understanding, and please don't hesitate to contact us with any further
        questions.\n\nBest regards,\nCustomer Support"
      },
      "finish_reason": "stop"
    }
  ],
  "usage": {
    "prompt_tokens": 96,
    "completion_tokens": 103,
    "total_tokens": 199
  }
}

```

Fields extracted by AIProviderResponseParser: - choices[0].message.content → reply text returned to user - usage.total_tokens → stored as token count

Error Response Examples

All providers return structured error responses. `AIProviderResponseParser` handles the following HTTP status codes:

401 / 403 — Authentication Failure

HTTP/1.1 401 Unauthorized
Content-Type: application/json

```
{
  "error": {
    "message": "Invalid API key provided.",
    "type": "invalid_request_error",
    "code": "invalid_api_key"
  }
}
```

Handling: Error detail logged server-side via `AIReplyLogger` at `LogLevel.ERROR`. Generic message returned to UI: *“AI provider authentication failed. Please verify your API key and Named Credential configuration.”*

429 — Rate Limit Exceeded

HTTP/1.1 429 Too Many Requests
Content-Type: application/json

```
{
  "error": {
    "message": "Rate limit reached for model gpt-4o.",
    "type": "requests",
    "code": "rate_limit_exceeded"
  }
}
```

Handling: Generic message to UI: *“AI provider rate limit exceeded. Please try again in a few moments.”*

404 — Model Not Found

HTTP/1.1 404 Not Found
Content-Type: application/json

```
{
  "error": {
    "message": "The model 'gpt-5-turbo' does not exist.",
    "type": "invalid_request_error",
    "code": "model_not_found"
  }
}
```

Handling: Generic message to UI: *“AI provider could not find the requested resource. Check that the Model name is correct and available for your account. Verify the Model value in your AI Reply Config.”*

Security Notes

Item	Detail
API key storage	Stored in Salesforce External Credential (AI_Email_Reply_ExtCred), encrypted at rest, never in Apex code
Authorization header	Injected automatically by Named Credential generateAuthorizationHeader; not set in Apex
TLS	All callouts use HTTPS (TLS 1.2+); enforced by Salesforce platform
Timeout	120,000ms (120 seconds) set on all requests
Callout allowlist	Enforced via Named Credential — no arbitrary endpoints callable
Error logging	Only e.getMessage() logged at LogLevel.ERROR; full stack traces never logged
UI error messages	Generic only — raw provider error details are never surfaced to end users
Data scope	Only email subject, body excerpt, tone, and template instructions are included in prompts; no Salesforce record IDs or user credentials